

## Loading the required libraries

```
In [1]: import pandas as pd
import numpy as np
from sklearn import preprocessing
import matplotlib.pyplot as plt
import seaborn as sb
from sklearn.impute import SimpleImputer
from sklearn import metrics
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from tabulate import tabulate
import warnings
warnings.filterwarnings("ignore")
```

## Loading the datasets

```
In [2]: X_train_df = pd.read_csv(r"D:\Monsoon Credittech\DS Test\DS Test\Training\X_train.csv")
X_train_df.head()
```

Out[2]:

	Unique_ID	C1	C2	C3	C4	C5	C6	C7	C8	N1	...	N26	N27	N28	N29
0	Candidate_5926	1	0	11	31	0	False	0	True	23.75	...	NaN	NaN	NaN	NaN
1	Candidate_48134	1	4	2	66	2	False	1	True	11.05	...	NaN	NaN	NaN	NaN
2	Candidate_51717	1	0	19	2	0	False	0	True	29.00	...	NaN	NaN	NaN	NaN
3	Candidate_26401	1	1	16	47	1	False	4	True	17.99	...	NaN	NaN	NaN	NaN
4	Candidate_34872	1	1	13	1	1	True	6	True	27.50	...	NaN	NaN	NaN	NaN

5 rows × 44 columns



```
In [3]: Y_train_df = pd.read_csv(r"D:\Monsoon Credittech\DS Test\DS Test\Training\y_train.csv")
Y_train_df.head()
```

Out[3]:

	Unique_ID	Dependent_Variable
0	Candidate_5926	1
1	Candidate_48134	0
2	Candidate_51717	1
3	Candidate_26401	0
4	Candidate_34872	0

```
In [4]: X_test_df = pd.read_csv(r"D:\Monsoon Credittech\DS Test\DS Test\Test\X_test.csv")
X_test_df.head(2)
```

Out[4]:

	Unique_ID	C1	C2	C3	C4	C5	C6	C7	C8	N1	...	N26	N27	N28	N29	I
0	Candidate_1602	1	0	0	23	0	True	0	True	18.00	...	NaN	NaN	NaN	NaN	NaN
1	Candidate_29650	1	0	2	4	2	True	2	True	16.75	...	NaN	NaN	NaN	NaN	NaN

2 rows × 44 columns



## Removing Unique\_ID feature from the training and test dataset

```
In [5]: Y_train = Y_train_df.drop(['Unique_ID'], axis=1)
X_train = X_train_df.drop(['Unique_ID'], axis=1)
X_test = X_test_df.drop(['Unique_ID'], axis=1)
print(Y_train.shape)
print(X_train.shape)
print(X_test.shape)
```

```
(33050, 1)
(33050, 43)
(11017, 43)
```

## Data Preprocessing

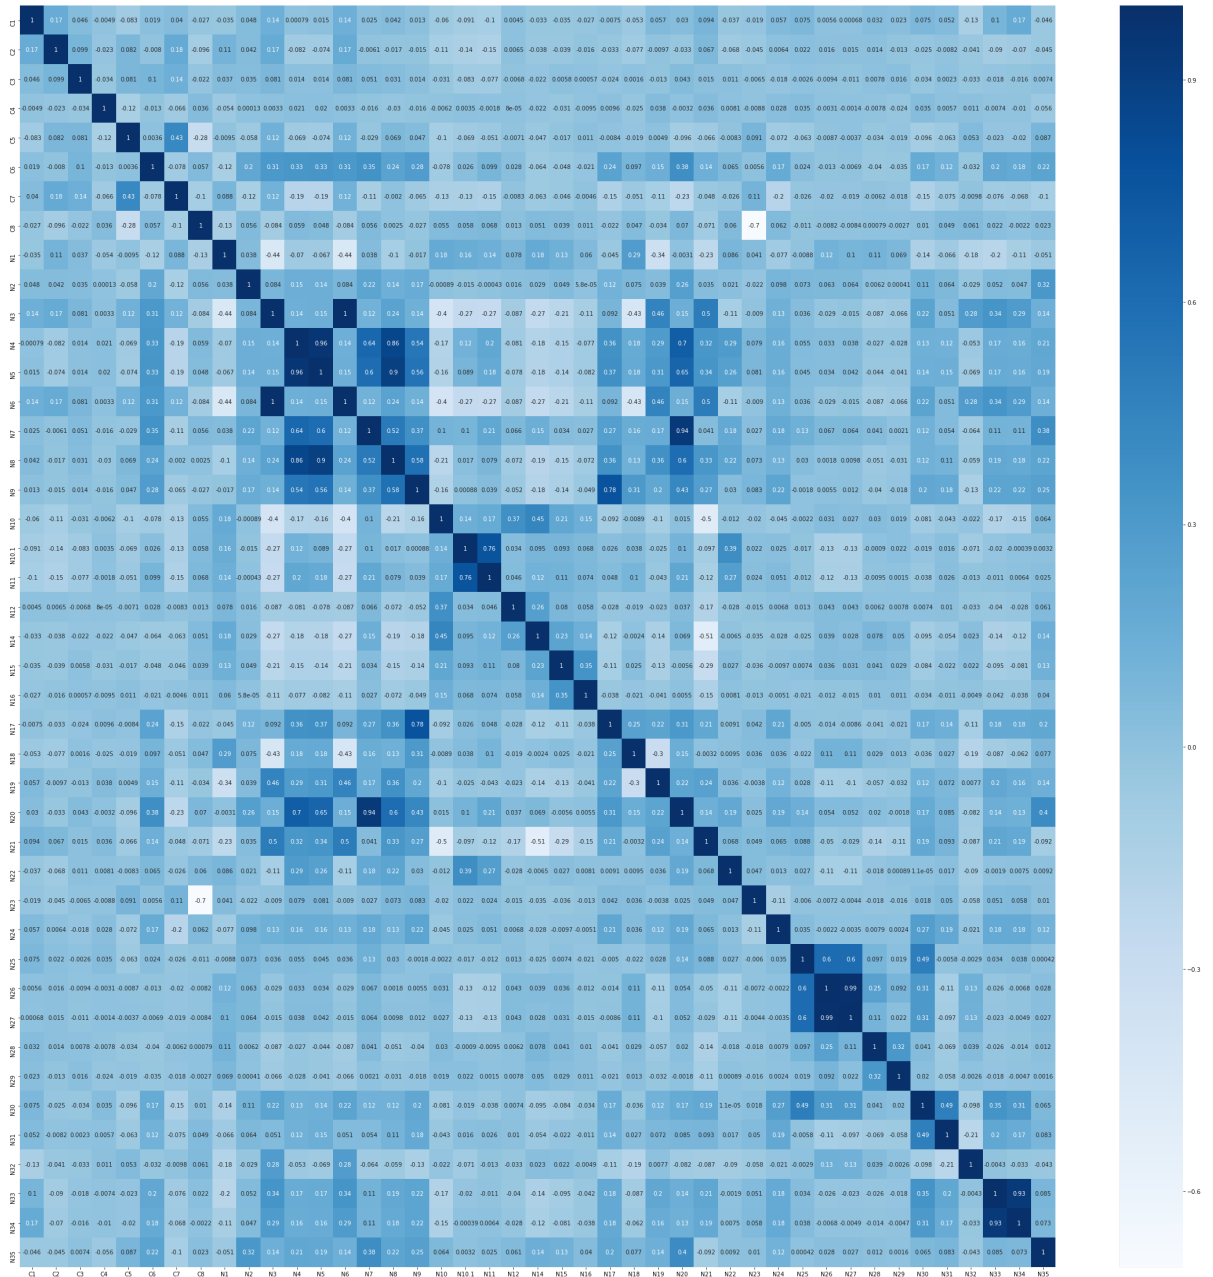
Converting boolean values (True and False) of columns "C6" and "C8" to 1 and 0 respectively

```
In [6]: X_train["C6"] = X_train["C6"].apply(lambda x:1 if x==True else 0)
X_train["C8"] = X_train["C8"].apply(lambda x:1 if x==True else 0)
```

```
In [7]: X_test["C6"] = X_test["C6"].apply(lambda x:1 if x==True else 0)
X_test["C8"] = X_test["C8"].apply(lambda x:1 if x==True else 0)
```

## Checking the correlation between features

```
In [8]: plt.figure(figsize = (40,40))
sb.heatmap(X_train.corr(), cmap="Blues", annot=True)
plt.show()
```



```
In [9]: corr_matrix = X_train.corr().abs()
```

```
In [10]: # Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(np.bool))
```

### Finding features having correlation greather than 90%

```
In [11]: to_drop = [column for column in upper.columns if any(upper[column] > 0.90)]
```

```
In [12]: to_drop
```

```
Out[12]: ['N5', 'N6', 'N20', 'N27', 'N34']
```

### Dropping features in the training dataset having correlation of greater than 0.9

```
In [13]: X_train.drop(to_drop, axis=1, inplace=True)
X_train.shape
```

```
Out[13]: (33050, 38)
```

### Dropping the same columns(columns that have been droppped from the train dataset having correlation of greater than 0.9) from the test dataset as well

```
In [14]: X_test.drop(to_drop, axis=1, inplace=True)
X_test.shape
```

```
Out[14]: (11017, 38)
```

### Handling missing values

```
In [15]: print("Columns having null data in the train dataset:\n",X_train.columns[X_train.isnull().any()].tolist())
```

Columns having null data in the train dataset:

```
['N2', 'N3', 'N4', 'N7', 'N10', 'N10.1', 'N11', 'N12', 'N14', 'N15', 'N16',
'N17', 'N18', 'N19', 'N21', 'N22', 'N23', 'N25', 'N26', 'N28', 'N29', 'N30',
'N31', 'N32', 'N35']
```

```
In [16]: print("Percentage of null data in the columns:")
print(((X_train[X_train.columns[X_train.isnull().any()]].isnull().sum()*100)/
X_train.shape[0])
```

Percentage of null data in the columns:

N2	13.954614
N3	1.104387
N4	13.936460
N7	1.291982
N10	1.291982
N10.1	1.291982
N11	2.166415
N12	13.960666
N14	1.839637
N15	1.291982
N16	13.936460
N17	13.936460
N18	13.936460
N19	13.830560
N21	13.830560
N22	13.830560
N23	7.521936
N25	81.025719
N26	81.025719
N28	81.025719
N29	81.025719
N30	81.025719
N31	81.025719
N32	81.180030
N35	1.291982

dtype: float64

### Finding columns having more than 80 percent of null values

```
In [17]: columns_with_null_gt_80 = X_train.columns[X_train.isna().sum()/len(X_train) >
0.8]
columns_with_null_gt_80
```

```
Out[17]: Index(['N25', 'N26', 'N28', 'N29', 'N30', 'N31', 'N32'], dtype='object')
```

### Dropping columns having more than 80 percent of null values

Dropping the columns having more than 80 percent of null data as the performance of the models after removing these columns is relatively same to the performance of models trained on the dataset wherein the null values in these columns are filled with their mean values.

Since we have no information about the column and its importance; we drop these columns as 80 percent is relatively high and also since there is no drop in performance.

```
In [18]: X_train.dropna(thresh=X_train.shape[0]*0.8,how='all',axis=1, inplace=True)
X_train.head()
```

Out[18]:

	C1	C2	C3	C4	C5	C6	C7	C8	N1	N2	...	N16	N17	N18	N19	N21	N2
0	1	0	11	31	0	0	0	1	23.75	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
1	1	4	2	66	2	0	1	1	11.05	22.0	...	0.0	1944.0	0.06	25856.0	0.88	1.0
2	1	0	19	2	0	0	0	1	29.00	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
3	1	1	16	47	1	0	4	1	17.99	1.0	...	0.0	8244.0	0.89	1006.0	1.00	0.0
4	1	1	13	1	1	1	6	1	27.50	206.0	...	0.0	57532.0	0.97	3398.0	0.96	0.0

5 rows × 31 columns



```
In [19]: X_train.shape
```

Out[19]: (33050, 31)

**Dropping the same columns(columns that have been dropped from the train dataset having more than 80 percent of null values) from the test dataset as well**

```
In [20]: X_test.drop(columns_with_null_gt_80, axis=1, inplace=True)
X_test.head()
```

Out[20]:

	C1	C2	C3	C4	C5	C6	C7	C8	N1	N2	...	N16	N17	N18	N19	N21	N2
0	1	0	0	23	0	1	0	1	18.00	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN
1	1	0	2	4	2	1	2	1	16.75	107.0	...	0.0	59435.0	0.83	12165.0	0.94	2.0
2	1	2	3	38	1	0	4	1	29.99	45.0	...	0.0	1996.0	0.79	504.0	0.70	2.0
3	1	1	28	20	2	0	2	1	17.70	20.0	...	0.0	9281.0	0.84	1428.0	0.77	0.0
4	1	1	15	1	3	0	5	0	28.00	2.0	...	0.0	13902.0	0.64	6324.0	0.92	2.0

5 rows × 31 columns



```
In [21]: print("Columns having null data in the train dataset:\n",X_train.columns[X_train.isnull().any()].tolist())
```

```
Columns having null data in the train dataset:
['N2', 'N3', 'N4', 'N7', 'N10', 'N10.1', 'N11', 'N12', 'N14', 'N15', 'N16',
'N17', 'N18', 'N19', 'N21', 'N22', 'N23', 'N35']
```

**Filling null values**

After trying various methods like dropping rows with null values; filling null values with mean, median, mode; and using interpolation, it was found that by filling null values with their mean there is an improvement in the performance. So, we fill the null values with mean.

If the class labels of the test dataset would have been available we could have tried filling the null values with mean of the data in the column having the same class label.

```
In [22]: imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
```

```
In [23]: X_train = pd.DataFrame(imputer.fit_transform(X_train), columns = X_train.columns)
```

```
In [24]: print("Columns having null data in the train dataset:\n",X_train.columns[X_train.isnull().any()].tolist())
```

```
Columns having null data in the train dataset:
[]
```

### Filling null values in the test dataset with their mean

```
In [25]: X_test = pd.DataFrame(imputer.fit_transform(X_test), columns = X_test.columns)
```


```
In [26]: print("Columns having null data in the test dataset:\n",X_test.columns[X_test.isnull().any()].tolist())
```

```
Columns having null data in the test dataset:
[]
```

### Checking the distribution of the class labels

```
In [27]: Y_train['Dependent_Variable'].value_counts().plot.pie(figsize=(6,4))
```

```
Out[27]: <AxesSubplot:ylabel='Dependent_Variable'>
```



We can see that the training data is highly imbalanced

**Using SMOTE (Synthetic Minority Oversampling Technique) to balance the class distribution by randomly increasing minority class examples by replicating them**

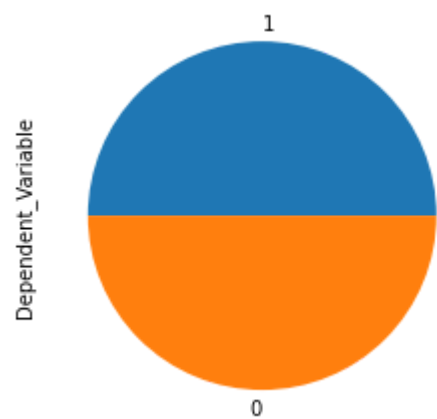
```
In [28]: smote = SMOTE(sampling_strategy='minority')
X_train, Y_train = smote.fit_resample(X_train, Y_train)
```

```
In [29]: print(X_train.shape)
print(Y_train.shape)
```

(45688, 31)
(45688, 1)

```
In [30]: Y_train['Dependent_Variable'].value_counts().plot.pie(figsize=(6,4))
```

Out[30]: <AxesSubplot:ylabel='Dependent\_Variable'>



```
In [31]: X_train.head(2)
```

Out[31]:

	C1	C2	C3	C4	C5	C6	C7	C8	N1	N2	...	N16	N17	
0	1.0	0.0	11.0	31.0	0.0	0.0	0.0	1.0	23.75	81.34271	...	0.023203	15820.726797	0.528
1	1.0	4.0	2.0	66.0	2.0	0.0	1.0	1.0	11.05	22.00000	...	0.000000	1944.000000	0.060

2 rows × 31 columns



```
In [32]: Y_train.head(2)
```

```
Out[32]:
```

	Dependent_Variable
0	1
1	0

```
In [33]: X_test.head(2)
```

```
Out[33]:
```

	C1	C2	C3	C4	C5	C6	C7	C8	N1	N2	...	N16	N17	
0	1.0	0.0	0.0	23.0	0.0	1.0	0.0	1.0	18.00	81.279492	...	0.024032	15937.365096	0.53
1	1.0	0.0	2.0	4.0	2.0	1.0	2.0	1.0	16.75	107.000000	...	0.000000	59435.000000	0.83

2 rows × 31 columns



## Preprocessing the numerical values

After trying both normalizing the numerical values and standardizing the numerical values, it was found that standardizing works better. So, we standardize the numerical values in the training and test dataset

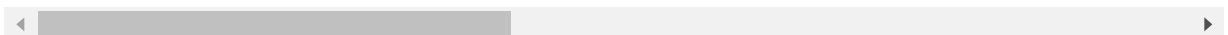
```
In [34]: scaler = preprocessing.StandardScaler()
```

```
In [35]: X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X_train.columns)
X_train.head()
```

```
Out[35]:
```

	C1	C2	C3	C4	C5	C6	C7	C8	
0	-0.302789	-0.719094	-0.158062	1.156846	-1.319561	-0.961579	-1.452604	0.30378	0.
1	-0.302789	0.512251	-0.946089	3.541748	0.216968	-0.961579	-0.986711	0.30378	-1
2	-0.302789	-0.719094	0.542405	-0.819216	-1.319561	-0.961579	-1.452604	0.30378	1.
3	-0.302789	-0.411258	0.279730	2.247087	-0.551297	-0.961579	0.410968	0.30378	-0
4	-0.302789	-0.411258	0.017055	-0.887356	-0.551297	1.120599	1.342754	0.30378	0.

5 rows × 31 columns

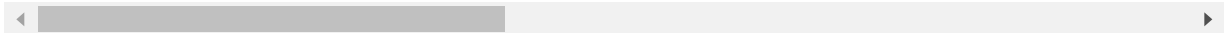


```
In [36]: X_test = pd.DataFrame(scaler.transform(X_test), columns = X_test.columns)
X_test.head()
```

Out[36]:

	C1	C2	C3	C4	C5	C6	C7	C8
0	-0.302789	-0.719094	-1.121206	0.611726	-1.319561	1.120599	-1.452604	0.303780
1	-0.302789	-0.719094	-0.946089	-0.682936	0.216968	1.120599	-0.520818	0.303780
2	-0.302789	-0.103421	-0.858530	1.633827	-0.551297	-0.961579	0.410968	0.303780
3	-0.302789	-0.411258	1.330432	0.407305	0.216968	-0.961579	-0.520818	0.303780
4	-0.302789	-0.411258	0.192171	-0.887356	0.985232	-0.961579	0.876861	-3.567258

5 rows × 31 columns



## Implementing Machine Learning Models

### Logistic Regression

```
In [37]: parameters = {"C": [10** -x for x in range(-3,3)]}

model = LogisticRegression(penalty='l2', random_state=0)
clf_log = GridSearchCV(model, parameters, cv=10, scoring='roc_auc', n_jobs=-1)
clf_log.fit(X_train, Y_train.values.ravel())

print("Best estimator=", clf_log.best_estimator_)
print("Best score =", clf_log.best_score_)

Best estimator= LogisticRegression(C=100, random_state=0)
Best score = 0.7502915727859788
```

```
In [38]: model_log = LogisticRegression(C= clf_log.best_params_['C'], penalty='l2', random_state=0)
model_log.fit(X_train, Y_train.values.ravel())

Y_pred_log = model_log.predict_proba(X_test)
```

### Decision Trees

```
In [39]: model_DT = DecisionTreeClassifier(random_state=0)
parameters={'criterion': ['gini','entropy','log_loss'] , 'min_samples_split':[4,8,10,12], 'max_features':['sqrt','log2','None']}
clf_DT = GridSearchCV(model_DT, parameters, cv=10, scoring='roc_auc', n_jobs=-1)
clf_DT.fit(X_train, Y_train.values.ravel())

print("Best estimator=", clf_DT.best_estimator_)
print("Best score=", clf_DT.best_score_)
```

```
Best estimator= DecisionTreeClassifier(criterion='entropy', max_features='sqrt',
                                     min_samples_split=12, random_state=0)
Best score= 0.7673087437186636
```

```
In [40]: model_DT = RandomForestClassifier(criterion = clf_DT.best_params_['criterion']
                                     , min_samples_split= clf_DT.best_params_['min_samples_split']
                                     , max_features = clf_DT.best_params_['max_features']
                                     , random_state=0)
model_DT.fit(X_train, Y_train.values.ravel())

Y_pred_DT = model_DT.predict_proba(X_test)
```

## Random Forest

```
In [43]: model_RF = RandomForestClassifier(random_state=0)
parameters={'criterion': ['gini','entropy','log_loss']
            , 'min_samples_split':[4,8,10,12]
            , 'n_estimators':[350, 400, 500]
            , 'max_features':['sqrt','log2','None']}
clf_RF = GridSearchCV(model_RF, parameters, cv=10, scoring='roc_auc', n_jobs=-1)
clf_RF.fit(X_train, Y_train.values.ravel())

print("Best estimator=", clf_RF.best_estimator_)
print("Best score=", clf_RF.best_score_)
```

```
Best estimator= RandomForestClassifier(criterion='entropy', max_features='sqrt',
                                     min_samples_split=4, n_estimators=500, random_state=0)
Best score= 0.9025488993051087
```

```
In [44]: model_RF = RandomForestClassifier(criterion = clf_RF.best_params_['criterion']
                                         , min_samples_split= clf_RF.best_params_['mi
n_samples_split']
                                         , max_features = clf_RF.best_params_['max_fe
atures']
                                         , n_estimators= clf_RF.best_params_['n_estim
ators']
                                         , random_state=0)
model_RF.fit(X_train, Y_train.values.ravel())

Y_pred_RF = model_RF.predict_proba(X_test)
```

## Gradient Boosting Decision Trees

```
In [46]: model_GBDT = GradientBoostingClassifier(random_state=0)
parameters = {'learning_rate':[0.1, 0.2, 0.3, 0.4]
              , 'min_samples_split':[4,8,10,12]
              , 'n_estimators':[200, 300, 400, 500]
              , 'max_features':['sqrt','log2','None']}
clf_GBDT = GridSearchCV(model_GBDT, parameters, cv=10, scoring='roc_auc', n_jo
bs=-1)
clf_GBDT.fit(X_train, Y_train.values.ravel())

print("Best estimator=", clf_GBDT.best_estimator_)
print("Best score=", clf_GBDT.best_score_)
```

```
Best estimator= GradientBoostingClassifier(learning_rate=0.2, max_features='s
qrt',
                                           min_samples_split=10, n_estimators=500,
                                           random_state=0)
Best score= 0.8931136232903331
```

```
In [47]: model_GBDT = GradientBoostingClassifier(learning_rate= clf_GBDT.best_params_['
learning_rate']
                                                , min_samples_split= clf_GBDT.best_par
ams_['min_samples_split']
                                                , n_estimators= clf_GBDT.best_params_['
n_estimators']
                                                , max_features= clf_GBDT.best_params_['
max_features']
                                                )
model_GBDT.fit(X_train, Y_train.values.ravel())

Y_pred_GBDT = model_GBDT.predict_proba(X_test)
```

## Summary

```
In [66]: # assigning data
table_data = [("Logistic Regression", clf_log.best_score_)
              ,("Decision Trees", clf_DT.best_score_)
              ,("Random Forest", clf_RF.best_score_)
              ,("Gradient Boosting Decision Trees", clf_GBDT.best_score_)
              ]

# creating header
table_head = ["Model", "Mean cross-validated roc-auc score calculated using GridSearchCV"]

# displaying the table
print(tabulate(table_data, headers=table_head, tablefmt="grid"))
```

```
+-----+-----+-----+
| Model | Mean cross-validated roc-auc score calculated using GridSearchCV |
+-----+-----+-----+
| Logistic Regression | 0.750292 |
+-----+-----+-----+
| Decision Trees | 0.767309 |
+-----+-----+-----+
| Random Forest | 0.902549 |
+-----+-----+-----+
| Gradient Boosting Decision Trees | 0.893114 |
+-----+-----+-----+
```

As we can see that the mean cross-validated roc-auc score using Random Forest is the highest, we will use Random Forest model to predict the class labels of the test dataset

## Predicting Dependent\_Variable using RandomForest model

```
In [71]: Class_1_Probability_RF = pd.DataFrame(Y_pred_RF, columns=model_DT.classes_)[1]
         .to_numpy()
         Class_1_Probability_RF
```

```
Out[71]: array([0.5236, 0.3793, 0.34496667, ..., 0.42028571, 0.3007,
                0.22995])
```

```
In [72]: result_df = pd.DataFrame(X_test_df['Unique_ID'].copy())
```

```
In [73]: result_df['Class_1_Probability'] = Class_1_Probability_RF  
result_df.head()
```

Out[73]:

	Unique_ID	Class_1_Probability
0	Candidate_1602	0.523600
1	Candidate_29650	0.379300
2	Candidate_31061	0.344967
3	Candidate_5768	0.203617
4	Candidate_27059	0.535033

## Dataframe to csv

```
In [77]: result_df.to_csv(r"D:\Monsoon Credittech\DS Test\DS Test\Test\final_predictions.csv", index = False)
```